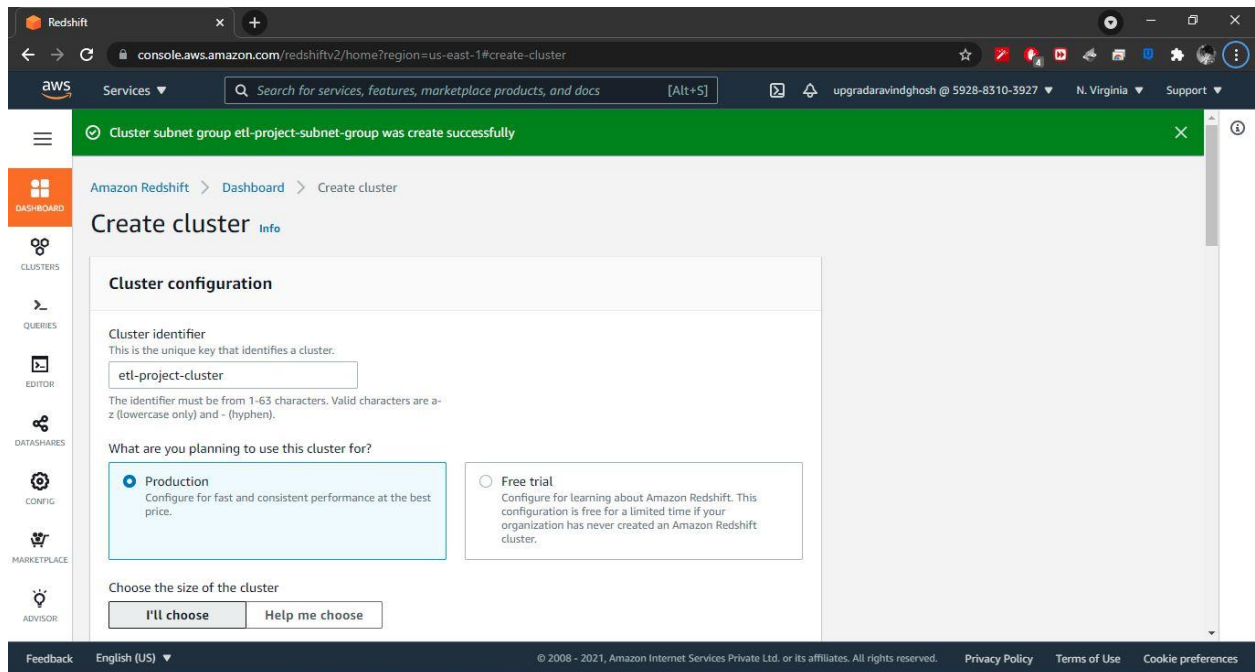
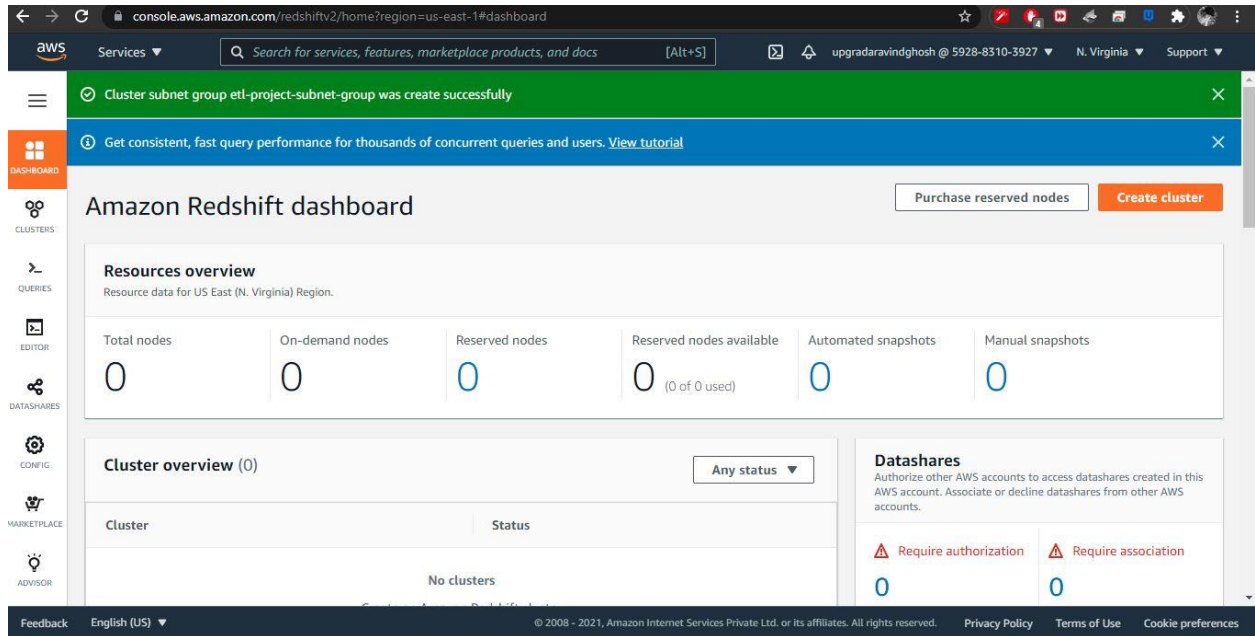


## Creation of a RedShift Cluster

Screenshots of the configuration of the RedShift cluster that you have created:



Redshift

console.aws.amazon.com/redshiftv2/home?region=us-east-1#create-cluster

Services

Search for services, features, marketplace products, and docs

[Alt+S]

upgradaravindghosh @ 5928-8310-3927

N. Virginia

Support

CONFIG

MARKETPLACE

ADVISOR

ALARMS

EVENTS

WHAT'S NEW

Production

Configure for fast and consistent performance at the best price.

Free trial

Configure for learning about Amazon Redshift. This configuration is free for a limited time if your organization has never created an Amazon Redshift cluster.

Choose the size of the cluster

I'll choose

Help me choose

Node type

Info

Choose a node type that meets your CPU, RAM, storage capacity, and drive type requirements.

dc2.large

Nodes

Enter the number of nodes that you need.

2

Range (1-32)

Configuration summary

Info

dc2.large | 2 nodes

\$360.00/month

Estimated on-demand compute price

320 GB

Total compressed storage

The total storage capacity for the

Feedback

English (US)

© 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

Redshift

console.aws.amazon.com/redshiftv2/home?region=us-east-1#create-cluster

Services

Search for services, features, marketplace products, and docs

[Alt+S]

upgradaravindghosh @ 5928-8310-3927

N. Virginia

Support

Sample data

Info

Load sample data

Load sample data to your Redshift cluster to start using the query editor to query data.

Database configurations

Admin user name

Enter a login ID for the admin user of your DB instance.

awsuser

The name must be 1-128 alphanumeric characters, and it can't be a reserved word

Auto generate password

Amazon Redshift can generate a password for you, or you can specify your own password.

Admin user password

Show password

Must be 8-64 characters long. Must contain at least one uppercase letter, one lowercase letter and one number. Can be any printable ASCII character except "/", """, or "@".

Cluster permissions

Feedback

English (US)

© 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

Redshift

console.aws.amazon.com/redshiftv2/home?region=us-east-1#create-cluster

Services

Search for services, features, marketplace products, and docs [Alt+S]

upgradaravindghosh @ 5928-8310-3927 N. Virginia Support

### Network and security

**Virtual private cloud (VPC)**  
This VPC defines the virtual networking environment for this cluster.

my\_vpc  
vpc-0674fb36dc09c54be

You can't change the VPC associated with this cluster after the cluster has been created. [Learn more](#)

**VPC security groups**  
This VPC security group defines which subnets and IP ranges the cluster can use in the VPC.

Choose one or more security groups

cloudera  
sg-0f38c3a73453d9d57

**Cluster subnet group**  
Choose the Amazon Redshift subnet group to launch the cluster in.

etl-project-subnet-group

**Availability Zone**  
Specify the Availability Zone that you want the cluster to be created in. Otherwise, Amazon Redshift chooses an Availability Zone for you.

No preference

Feedback English (US) © 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences

Redshift

console.aws.amazon.com/redshiftv2/home?region=us-east-1#create-cluster

Services

Search for services, features, marketplace products, and docs [Alt+S]

upgradaravindghosh @ 5928-8310-3927 N. Virginia Support

### Database configurations

**Database name**  
Specify a database name to create an additional database.

dev

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

**Database port**  
Port number where the database accepts inbound connections. You can't change the port after the cluster has been created.

5439

The port must be numeric (1150-65535).

**Parameter groups**  
Defines database parameter and query queues for all the databases.

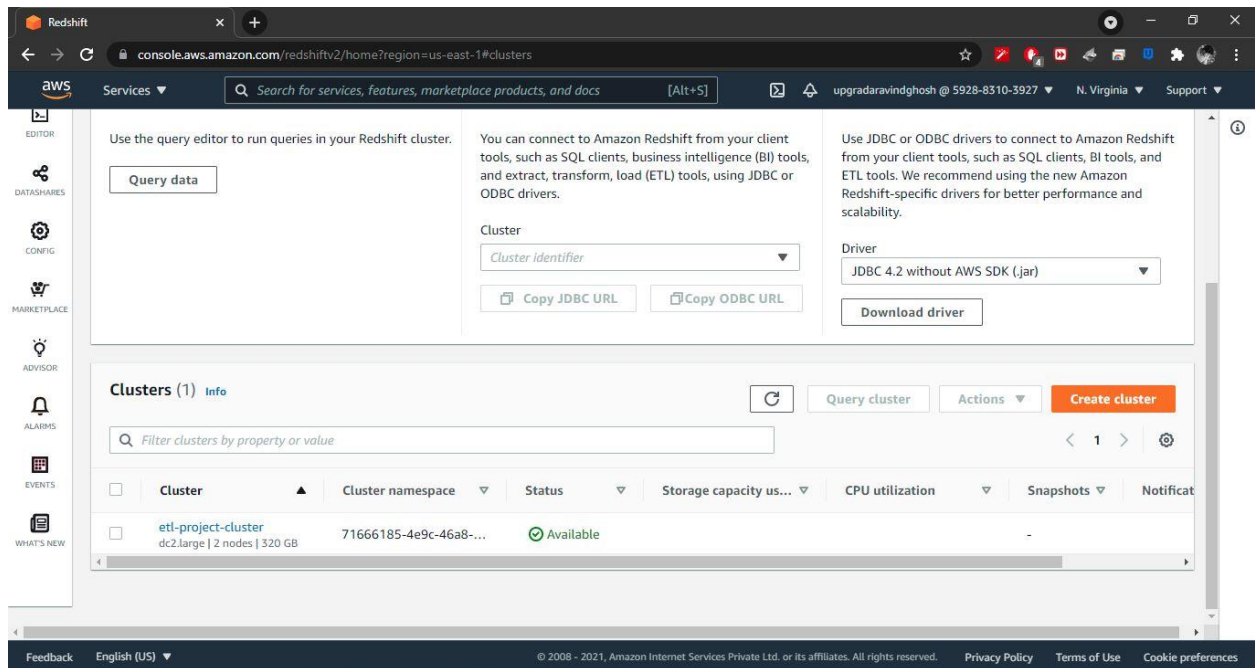
**Encryption**  
Encrypt all data on your cluster.

☒ Disabled

☐ Use AWS Key Management Service (AWS KMS)

☐ Use a hardware security module (HSM)

Feedback English (US) © 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences



Use the query editor to run queries in your Redshift cluster.

You can connect to Amazon Redshift from your client tools, such as SQL clients, business intelligence (BI) tools, and extract, transform, load (ETL) tools, using JDBC or ODBC drivers.

Use JDBC or ODBC drivers to connect to Amazon Redshift from your client tools, such as SQL clients, BI tools, and ETL tools. We recommend using the new Amazon Redshift-specific drivers for better performance and scalability.

Cluster:

Driver:

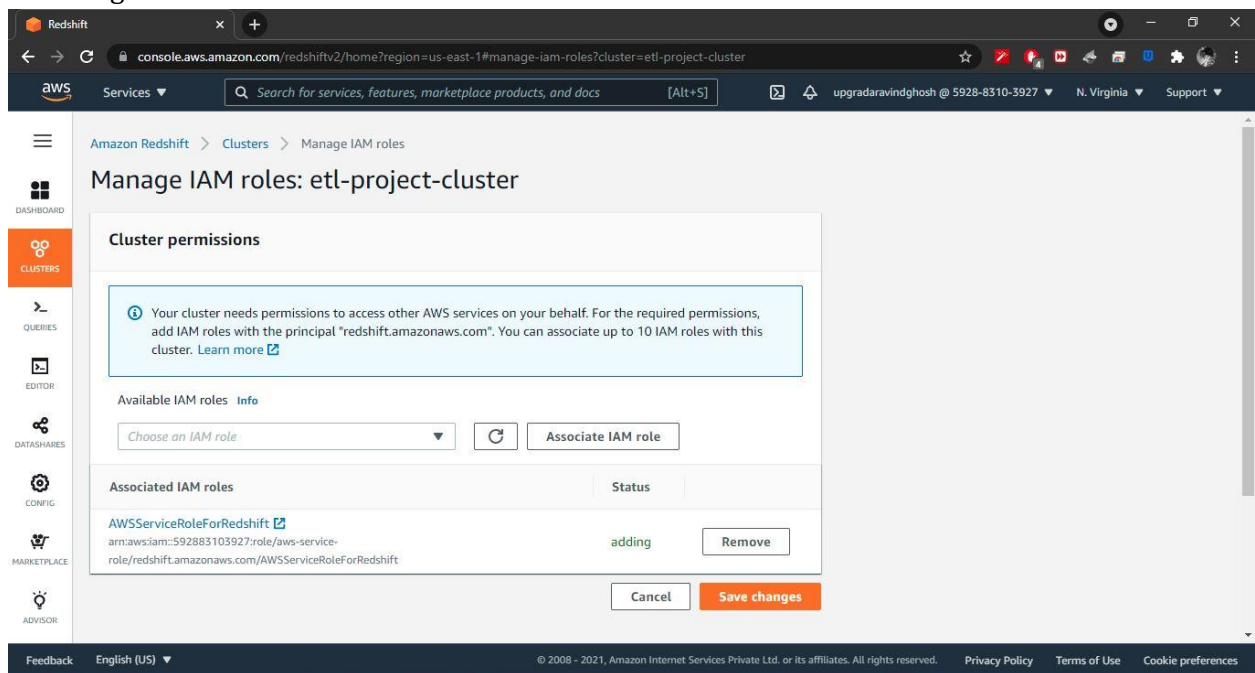
[Copy JDBC URL](#) [Copy ODBC URL](#) [Download driver](#)

**Clusters (1) Info**

Cluster	Cluster namespace	Status	Storage capacity us...	CPU utilization	Snapshots	Notificat
<input type="checkbox"/> etl-project-cluster dc2.large   2 nodes   320 GB	71666185-4e9c-46a8-...	Available				

Feedback English (US) © 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences

## Attaching IAM Roles:



Amazon Redshift > Clusters > Manage IAM roles

### Manage IAM roles: etl-project-cluster

**Cluster permissions**

**Available IAM roles** [Info](#)

[Associate IAM role](#)

Associated IAM roles	Status
<a href="#">AWSServiceRoleForRedshift</a> arn:aws:iam::592883103927:role/aws-service-role/redshift.amazonaws.com/AWSServiceRoleForRedshift	adding <a href="#">Remove</a>

[Cancel](#) [Save changes](#)

Feedback English (US) © 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences

The screenshot shows the AWS Redshift console interface. At the top, the browser address bar displays the URL: `console.aws.amazon.com/redshiftv2/home?region=us-east-1#cluster-details?cluster=etl-project-cluster&tab=properties`. The console header includes the AWS logo, a 'Services' dropdown, a search bar, and user information for 'upgradaravindghosh' in the 'N. Virginia' region.

The main content area is titled 'etl-project-cluster' and displays the following properties:

- Subnet:** etl-project-subnet-group
- Enhanced VPC routing:** Disabled
- Endpoint URL:** sg-0f38c3a73453d9d57
- Endpoint:** Disabled

Below the properties, the 'Cluster permissions (1)' section is visible. It includes a description: 'Your cluster needs permissions to access other AWS services on your behalf. For the required permissions, add IAM roles with the principal "redshift.amazonaws.com". You can associate up to 10 IAM roles with this cluster. [Learn more](#)'.

A 'Manage IAM roles' button is present. Below it, a table lists the associated IAM roles:

Associated IAM roles	Status	Amazon Resource Name (ARN)
<a href="#">AWSServiceRoleForRedshift</a>	in-sync	arn:aws:iam::592883103927:role/aws-service-role/redshift.amazonaws.com/AWSServiceRoleForRedshift

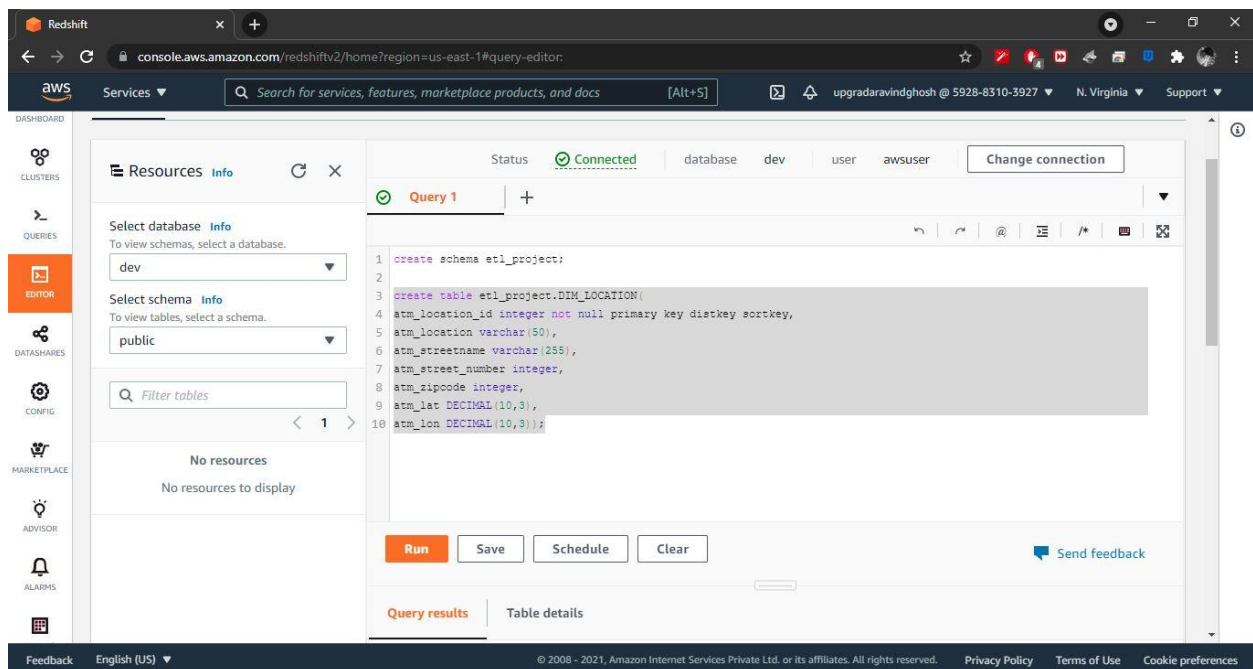
The 'Granted accounts (0)' section is also visible, with a description: 'VPcs in other accounts that are allowed to access this cluster. [Learn more](#)'. It includes buttons for 'Edit', 'Revoke', and 'Grant access', and a search bar for finding accounts or VPCs.

## Setting up a database in the RedShift cluster and running queries to create the dimension and fact tables

### Queries to create the various dimension and fact tables with appropriate primary and foreign keys:

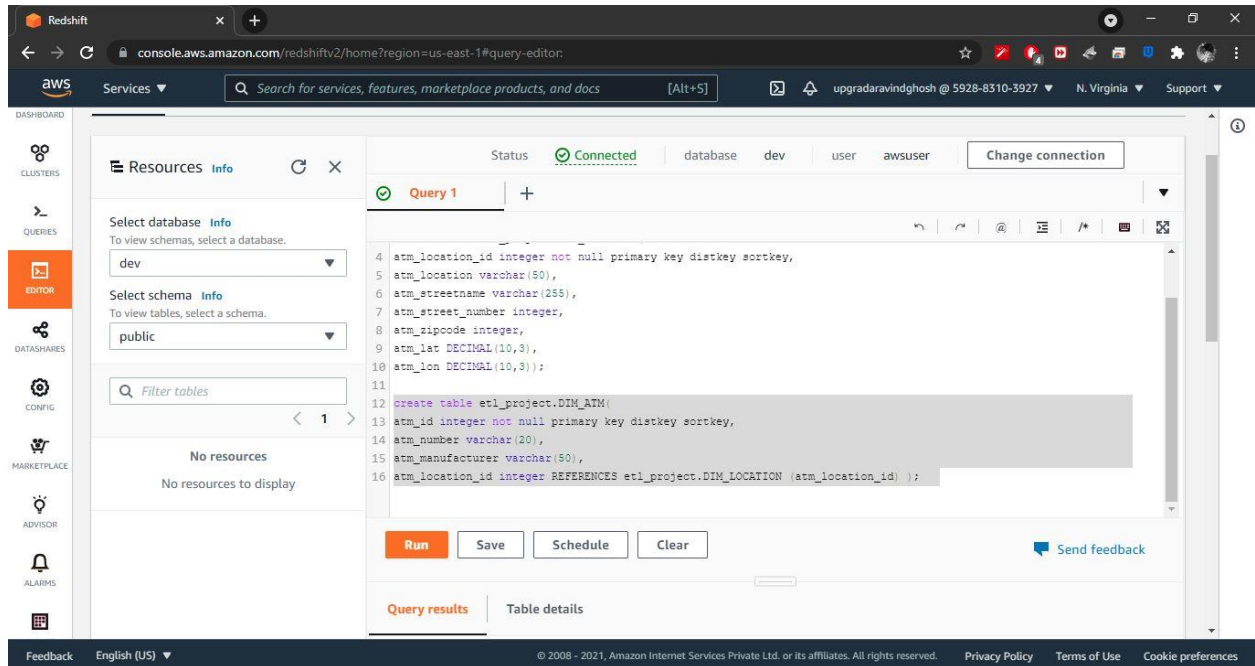
create schema etl\_project;

```
create table etl_project.DIM_LOCATION(  
  atm_location_id integer not null primary key distkey sortkey,  
  atm_location varchar(50),  
  atm_streetname varchar(255),  
  atm_street_number integer,  
  atm_zipcode integer,  
  atm_lat DECIMAL(10,3),  
  atm_lon DECIMAL(10,3));
```

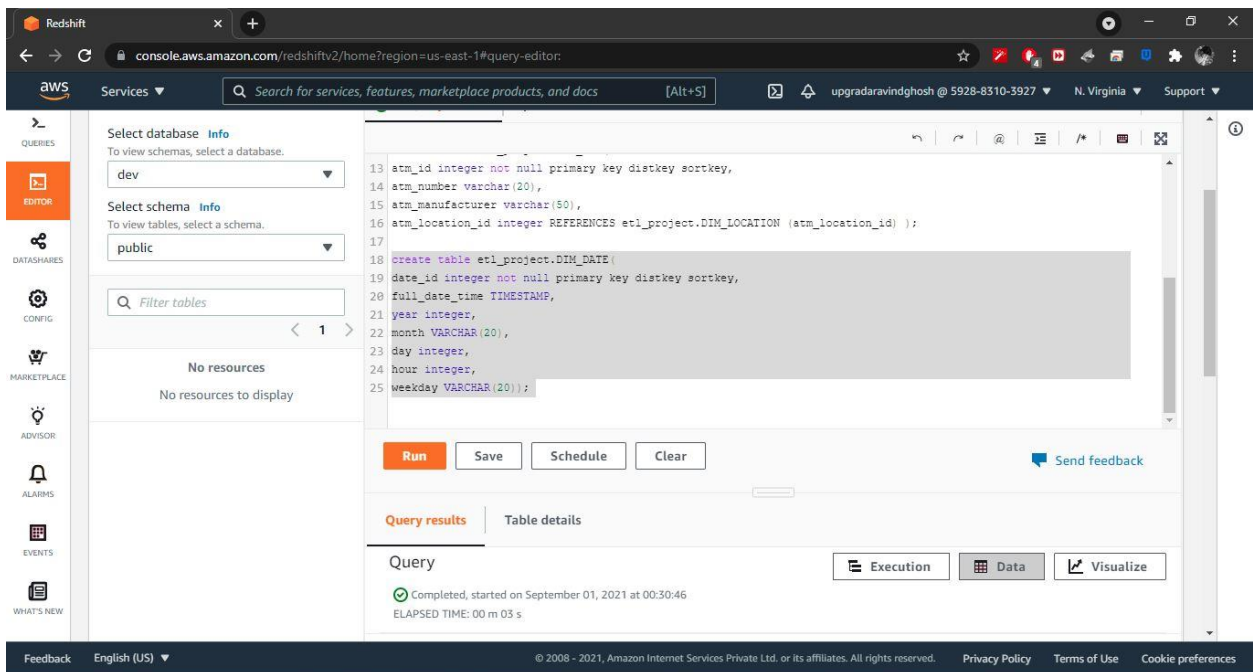




```
create table etl_project.DIM_ATM(  
atm_id integer not null primary key distkey sortkey,  
atm_number varchar(20),  
atm_manufacturer varchar(50),  
atm_location_id integer REFERENCES etl_project.DIM_LOCATION (atm_location_id) );
```



```
create table etl_project.DIM_DATE(  
date_id integer not null primary key distkey sortkey,  
full_date_time TIMESTAMP,  
year integer,  
month VARCHAR(20),  
day integer,  
hour integer,  
weekday VARCHAR(20));
```



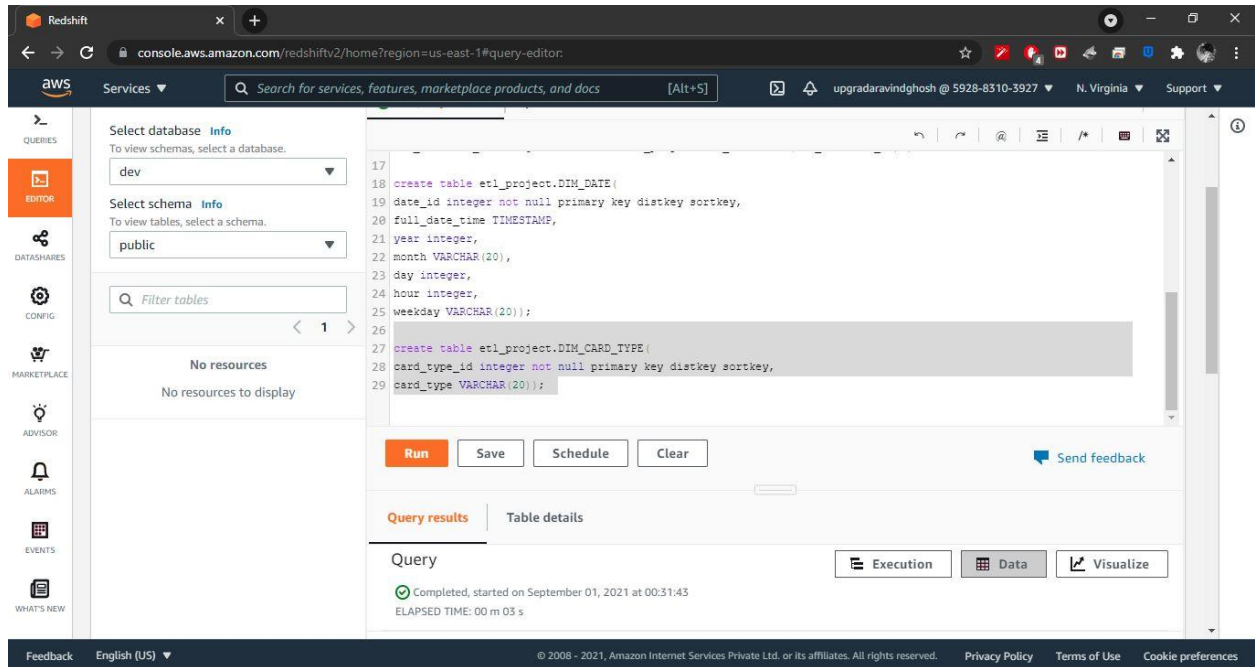
The screenshot shows the AWS Redshift console interface. The left sidebar contains navigation options: QUERIES, EDITOR (selected), DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area is divided into two panels. The left panel shows the 'Select database' dropdown set to 'dev' and the 'Select schema' dropdown set to 'public'. Below these, there is a 'Filter tables' search bar and a message 'No resources' with 'No resources to display'. The right panel displays a SQL query in a code editor. The query is as follows:

```
13 atm_id integer not null primary key distkey sortkey,  
14 atm_number varchar(20),  
15 atm_manufacturer varchar(50),  
16 atm_location_id integer REFERENCES etl_project.DIM_LOCATION (atm_location_id) );  
17  
18 Create table etl_project.DIM_DATE(  
19 date_id integer not null primary key distkey sortkey,  
20 full_date_time TIMESTAMP,  
21 year integer,  
22 month VARCHAR(20),  
23 day integer,  
24 hour integer,  
25 weekday VARCHAR(20));
```

Below the code editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. A 'Send feedback' link is also present. The 'Query results' tab is selected, showing a status of 'Completed, started on September 01, 2021 at 00:30:46' and 'ELAPSED TIME: 00 m 03 s'. There are also buttons for 'Execution', 'Data', and 'Visualize'.



```
create table etl_project.DIM_CARD_TYPE(  
card_type_id integer not null primary key distkey sortkey,  
card_type VARCHAR(20));
```

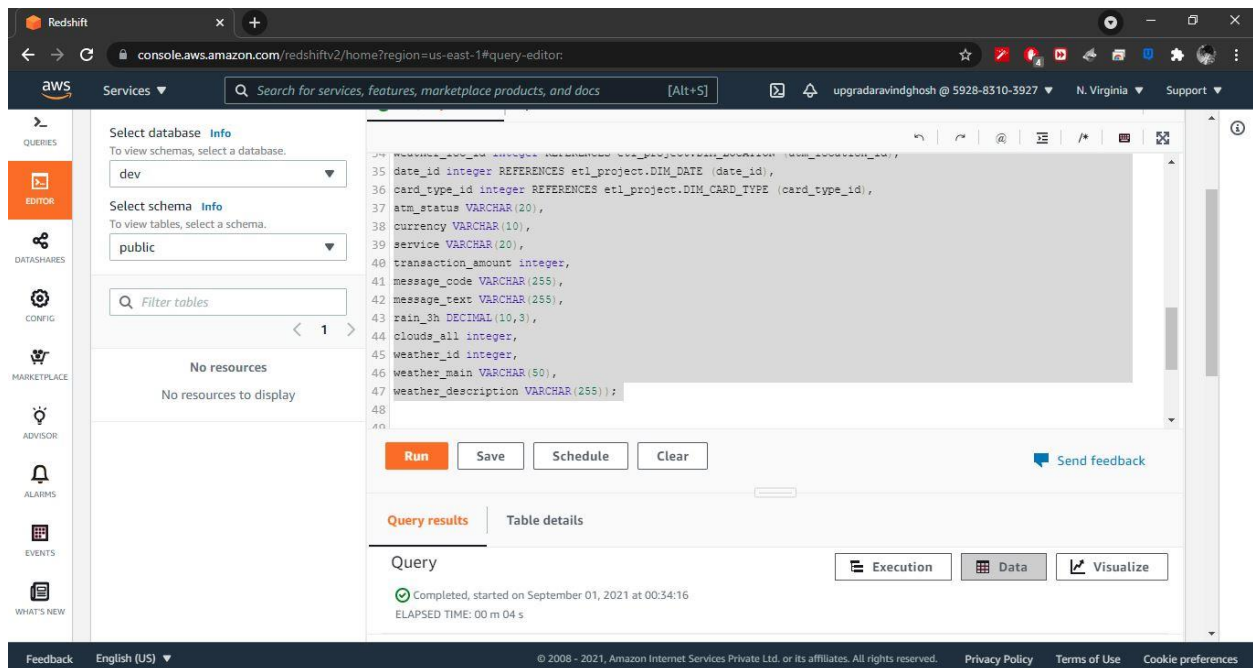


The screenshot shows the AWS Redshift console interface. On the left, there's a sidebar with navigation options like QUERIES, EDITOR, DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area is divided into two panes. The left pane shows the 'Select database' dropdown set to 'dev' and the 'Select schema' dropdown set to 'public'. Below this, it says 'No resources' and 'No resources to display'. The right pane contains a SQL query editor with the following code:

```
17  
18 create table etl_project.DIM_DATE(  
19 date_id integer not null primary key distkey sortkey,  
20 full_date_time TIMESTAMP,  
21 year integer,  
22 month VARCHAR(20),  
23 day integer,  
24 hour integer,  
25 weekday VARCHAR(20));  
26  
27 create table etl_project.DIM_CARD_TYPE(  
28 card_type_id integer not null primary key distkey sortkey,  
29 card_type VARCHAR(20));
```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. To the right of these buttons is a 'Send feedback' link. Below the buttons, there are tabs for 'Query results' and 'Table details'. The 'Query results' tab is active, showing a status of 'Completed, started on September 01, 2021 at 00:31:43' and 'ELAPSED TIME: 00 m 03 s'. At the bottom of the console, there's a footer with 'Feedback', 'English (US)', and copyright information: '© 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.' along with links for 'Privacy Policy', 'Terms of Use', and 'Cookie preferences'.

```
create table etl_project.FACT_ATM_TRANS(
trans_id BIGINT not null primary key distkey sortkey,
atm_id integer REFERENCES etl_project.DIM_ATM (atm_id),
weather_loc_id integer REFERENCES etl_project.DIM_LOCATION (atm_location_id),
date_id integer REFERENCES etl_project.DIM_DATE (date_id),
card_type_id integer REFERENCES etl_project.DIM_CARD_TYPE (card_type_id),
atm_status VARCHAR(20),
currency VARCHAR(10),
service VARCHAR(20),
transaction_amount integer,
message_code VARCHAR(255),
message_text VARCHAR(255),
rain_3h DECIMAL(10,3),
clouds_all integer,
weather_id integer,
weather_main VARCHAR(50),
weather_description VARCHAR(255));
```



## Loading data into a RedShift cluster from Amazon S3 bucket

### Queries to copy the data from S3 buckets to the RedShift cluster in the appropriate tables

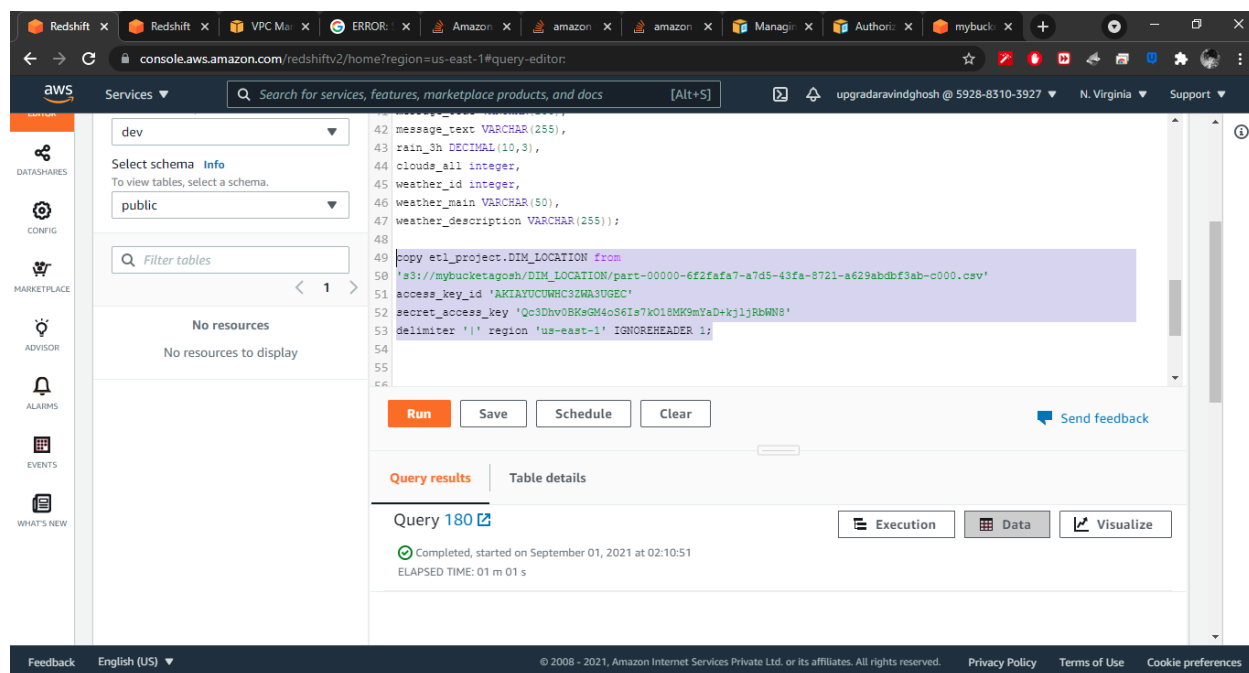
copy etl\_project.DIM\_LOCATION from

's3://mybucketagosh/DIM\_LOCATION/part-00000-6f2fafa7-a7d5-43fa-8721-a629abdbf3ab-c000.csv'

access\_key\_id 'AKIAYUCUWHC3ZWA3UGEC'

secret\_access\_key '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'

delimiter '|' region 'us-east-1' IGNOREHEADER 1;



The screenshot shows the AWS RedShift console interface. The left sidebar contains navigation options like DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area displays a SQL query editor with the following code:

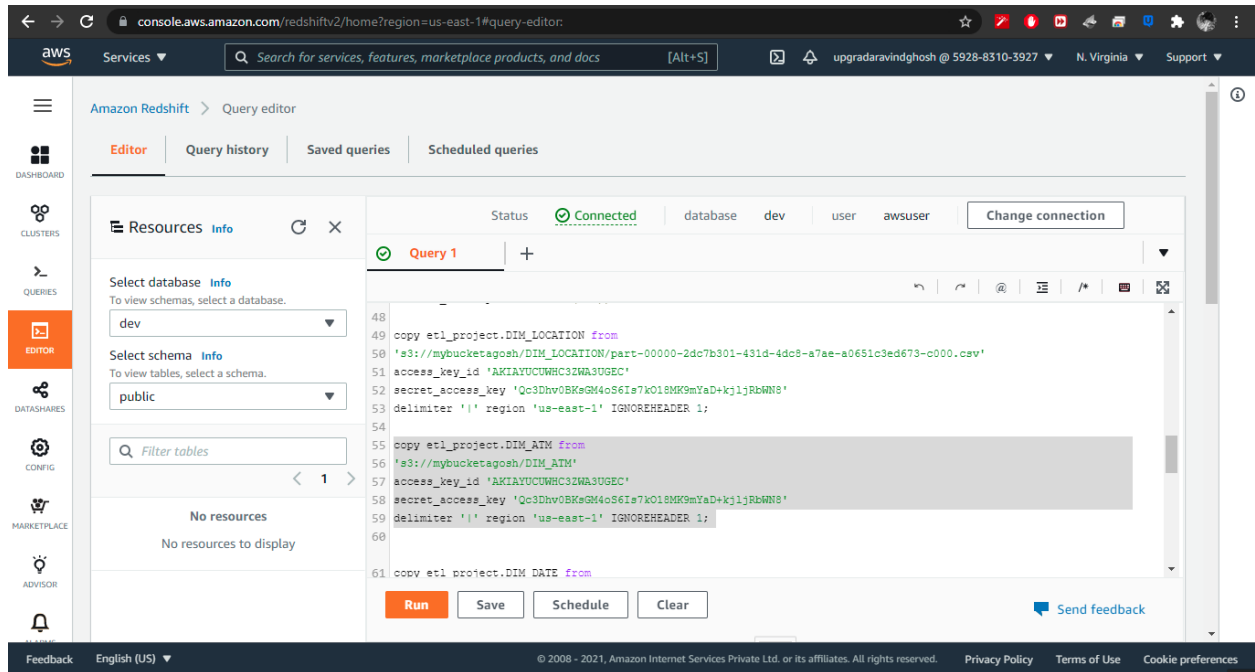
```

42 message_text VARCHAR(255),
43 rain_3h DECIMAL(10,3),
44 clouds_all integer,
45 weather_id integer,
46 weather_main VARCHAR(50),
47 weather_description VARCHAR(255));
48
49 copy etl_project.DIM_LOCATION from
50 's3://mybucketagosh/DIM_LOCATION/part-00000-6f2fafa7-a7d5-43fa-8721-a629abdbf3ab-c000.csv'
51 access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
52 secret_access_key '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'
53 delimiter '|' region 'us-east-1' IGNOREHEADER 1;
54
55
56

```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted. To the right of the buttons is a 'Send feedback' link. Below the buttons, the 'Query results' tab is selected, showing a status of 'Completed, started on September 01, 2021 at 02:10:51' and 'ELAPSED TIME: 01 m 01 s'. The 'Table details' tab is also visible. At the bottom of the console, there is a footer with 'Feedback', 'English (US)', and copyright information for Amazon Internet Services Private Ltd.

copy etl\_project.DIM\_ATM from  
's3://mybucketagosh/DIM\_ATM'  
access\_key\_id 'AKIAYUCUWHC3ZWA3UGEC'  
secret\_access\_key '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'  
delimiter '|' region 'us-east-1' IGNOREHEADER 1;



The screenshot shows the Amazon Redshift Query Editor interface. The left sidebar contains navigation options: DASHBOARD, CLUSTERS, QUERIES, EDITOR (selected), DATASHARES, CONFIG, MARKETPLACE, and ADVISOR. The main area is titled 'Amazon Redshift > Query editor' and has tabs for Editor, Query history, Saved queries, and Scheduled queries. The 'Editor' tab is active, showing a SQL query in a text editor. The query is as follows:

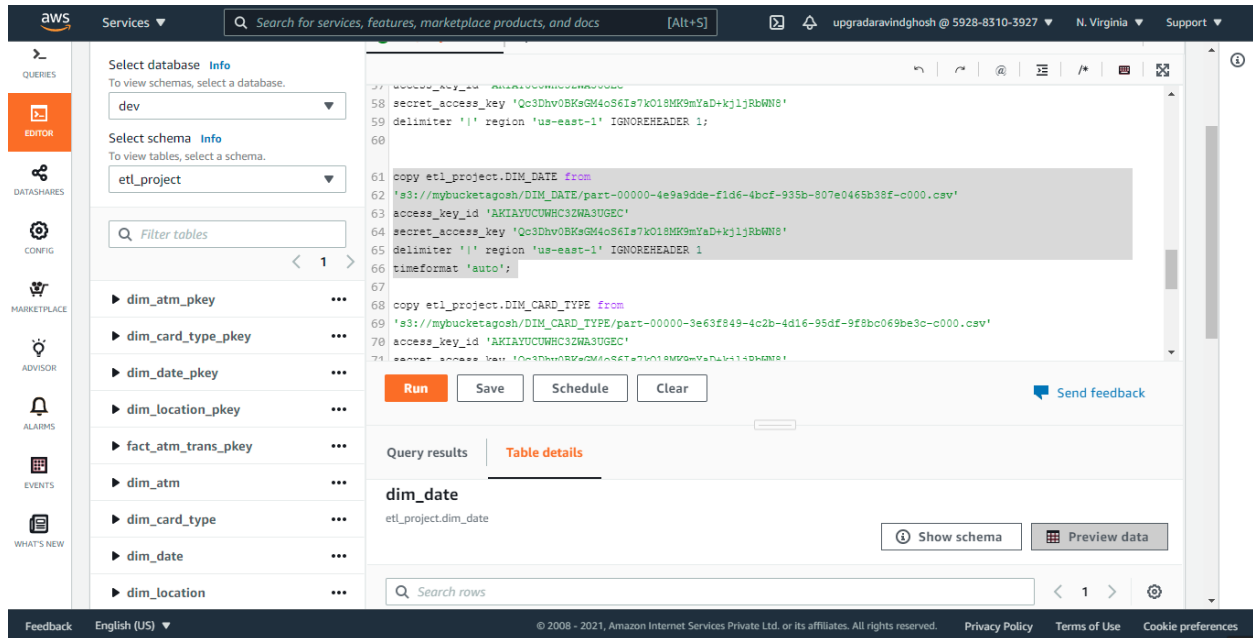
```

48 copy etl_project.DIM_LOCATION from
49 's3://mybucketagosh/DIM_LOCATION/part-00000-2dc7b301-431d-4dc8-a7ae-a0651c3ed673-c000.csv'
50 access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
51 secret_access_key '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'
52 delimiter '|' region 'us-east-1' IGNOREHEADER 1;
53
54
55 copy etl_project.DIM_ATM from
56 's3://mybucketagosh/DIM_ATM'
57 access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
58 secret_access_key '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'
59 delimiter '|' region 'us-east-1' IGNOREHEADER 1;
60
61 copy etl_project.DIM DATE from

```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted in orange. The status bar at the bottom indicates 'Connected' and shows the current connection details: database, dev, user, awsuser. The footer of the page contains copyright information: '© 2008 - 2021, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.' and links to 'Privacy Policy', 'Terms of Use', and 'Cookie preferences'.

copy etl\_project.DIM\_DATE from  
's3://mybucketagosh/DIM\_DATE/part-00000-4e9a9dde-f1d6-4bcf-935b-807e0465b38f-c000.csv'  
access\_key\_id 'AKIAYUCUWHC3ZWA3UGEC'  
secret\_access\_key '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'  
delimiter '|' region 'us-east-1' IGNOREHEADER 1  
timeformat 'auto';

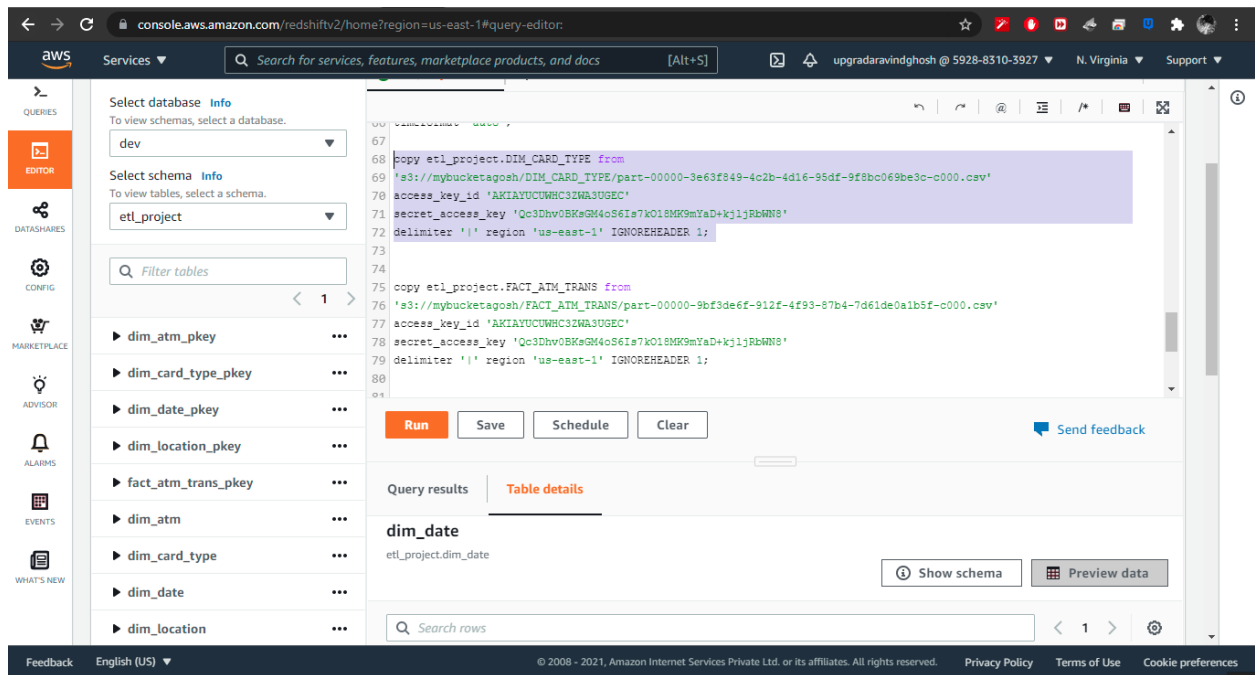


The screenshot shows the AWS Glue console interface. On the left, there's a sidebar with navigation options like QUERIES, EDITOR, DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area is divided into two panes. The left pane shows the 'dev' database and 'etl\_project' schema, with a list of tables including 'dim\_atm\_pkey', 'dim\_card\_type\_pkey', 'dim\_date\_pkey', 'dim\_location\_pkey', 'fact\_atm\_trans\_pkey', 'dim\_atm', 'dim\_card\_type', 'dim\_date', and 'dim\_location'. The right pane shows a SQL query editor with the following code:

```
57 secret_access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
58 secret_access_key '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'
59 delimiter '|' region 'us-east-1' IGNOREHEADER 1;
60
61 copy etl_project.DIM_DATE from
62 's3://mybucketagosh/DIM_DATE/part-00000-4e9a9dde-f1d6-4bcf-935b-807e0465b38f-c000.csv'
63 access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
64 secret_access_key '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'
65 delimiter '|' region 'us-east-1' IGNOREHEADER 1
66 timeformat 'auto';
67
68 copy etl_project.DIM_CARD_TYPE from
69 's3://mybucketagosh/DIM_CARD_TYPE/part-00000-3e63f849-4c2b-4d16-95df-9f8bc069be3c-c000.csv'
70 access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
71 secret_access_key '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';
```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. To the right of these buttons is a 'Send feedback' link. Below the query editor, there's a tabbed interface with 'Query results' and 'Table details'. The 'Table details' tab is active, showing the 'dim\_date' table. It lists the table name 'etl\_project.dim\_date' and has buttons for 'Show schema' and 'Preview data'. At the bottom, there's a search bar for rows and a pagination control showing '1'.

copy etl\_project.DIM\_CARD\_TYPE from  
's3://mybucketagosh/DIM\_CARD\_TYPE/part-00000-3e63f849-4c2b-4d16-95df-9f8bc069be3c-c000.csv'  
access\_key\_id 'AKIAYUCUWHC3ZWA3UGEC'  
secret\_access\_key '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'  
delimiter '|' region 'us-east-1' IGNOREHEADER 1;



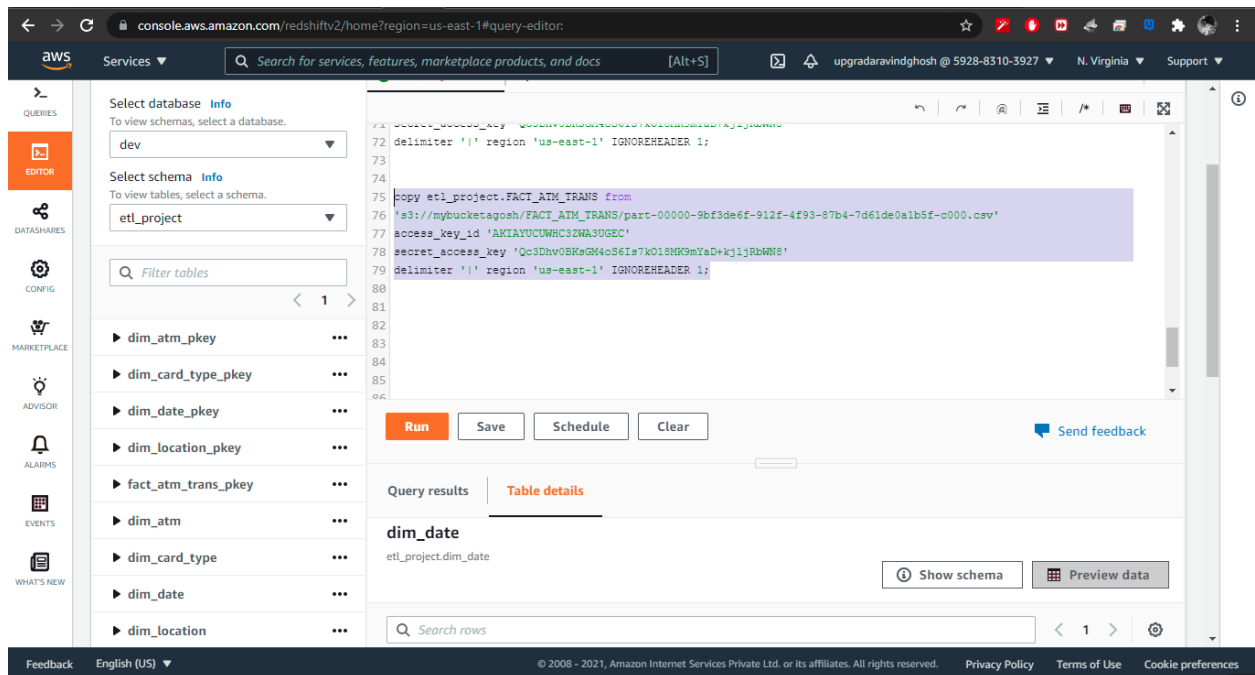
The screenshot shows the AWS Redshift console interface. On the left, there's a sidebar with navigation options like QUERIES, EDITOR, DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area is divided into two panes. The left pane shows the 'dev' database and 'etl\_project' schema, with a list of tables including dim\_atm\_pkey, dim\_card\_type\_pkey, dim\_date\_pkey, dim\_location\_pkey, fact\_atm\_trans\_pkey, dim\_atm, dim\_card\_type, dim\_date, and dim\_location. The right pane shows a SQL query editor with the following code:

```
copy etl_project.DIM_CARD_TYPE from
's3://mybucketagosh/DIM_CARD_TYPE/part-00000-3e63f849-4c2b-4d16-95df-9f8bc069be3c-c000.csv'
access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
secret_access_key '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'
delimiter '|' region 'us-east-1' IGNOREHEADER 1;
```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. To the right of these buttons is a 'Send feedback' link. Below the query editor, there's a section for 'Query results' and 'Table details'. The 'Table details' section shows the table 'dim\_date' with the schema 'etl\_project.dim\_date'. There are buttons for 'Show schema' and 'Preview data'. At the bottom, there's a search bar for rows and a pagination control showing '1'.



copy etl\_project.FACT\_ATM\_TRANS from  
 's3://mybucketagosh/FACT\_ATM\_TRANS/part-00000-9bf3de6f-912f-4f93-87b4-7d61de0a1b5f-c000.csv'  
 access\_key\_id 'AKIAYUCUWHC3ZWA3UGEC'  
 secret\_access\_key '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'  
 delimiter '|' region 'us-east-1' IGNOREHEADER 1;



The screenshot shows the AWS Redshift console query editor interface. The left sidebar contains navigation options like QUERIES, EDITOR, DATASHARES, CONFIG, MARKETPLACE, ADVISOR, ALARMS, EVENTS, and WHAT'S NEW. The main area is divided into two panes. The left pane shows the 'dev' database and 'etl\_project' schema, with a list of tables including dim\_atm\_pkey, dim\_card\_type\_pkey, dim\_date\_pkey, dim\_location\_pkey, fact\_atm\_trans\_pkey, dim\_atm, dim\_card\_type, dim\_date, and dim\_location. The right pane displays a SQL query for copying data from an S3 bucket to the 'dim\_date' table in the 'etl\_project' schema. The query is as follows:

```
copy etl_project.FACT_ATM_TRANS from
's3://mybucketagosh/FACT_ATM_TRANS/part-00000-9bf3de6f-912f-4f93-87b4-7d61de0a1b5f-c000.csv'
access_key_id 'AKIAYUCUWHC3ZWA3UGEC'
secret_access_key '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'
delimiter '|' region 'us-east-1' IGNOREHEADER 1;
```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted. Below these buttons, the 'Query results' tab is active, showing the 'dim\_date' table details. The table is located in the 'etl\_project' database. There are buttons for 'Show schema' and 'Preview data'. At the bottom of the console, there is a footer with 'Feedback', 'English (US)', and copyright information for Amazon Internet Services Private Ltd. (2008-2021).

**SELECT COUNT(\*) FROM etl\_project.DIM\_LOCATION;**

```

81 SELECT COUNT(*) FROM etl_project.DIM_LOCATION;
82 SELECT COUNT(*) FROM etl_project.DIM_ATM;
83 SELECT COUNT(*) FROM etl_project.DIM_DATE;
84 SELECT COUNT(*) FROM etl project.DIM CARD TYPE;

```

**Run** **Save** **Schedule** **Clear** [Send feedback](#)

**Query results** | **Table details**

Query **1381** [🔗](#) **Execution** **Data** **Visualize**

✓ Completed, started on September 08, 2021 at 22:49:27  
ELAPSED TIME: 00 m 02 s

**Rows returned (1)** **Export** ▼

< 1 > ⚙️

count
109

**SELECT COUNT(\*) FROM etl\_project.DIM\_ATM;**

```

81 SELECT COUNT(*) FROM etl_project.DIM_LOCATION;
82 SELECT COUNT(*) FROM etl_project.DIM_ATM;
83 SELECT COUNT(*) FROM etl_project.DIM_DATE;
84 SELECT COUNT(*) FROM etl_project.DIM_CARD_TYPE;
85 SELECT COUNT(*) FROM etl_project.FACT_ATM_TRANS;
86

```

**Run** **Save** **Schedule** **Clear** [Send feedback](#)

**Query results** | **Table details**

Query **1531** [🔗](#) **Execution** **Data** **Visualize**

✓ Completed, started on September 08, 2021 at 23:00:16  
ELAPSED TIME: 00 m 02 s

**Rows returned (1)** **Export** ▼

< 1 > ⚙️

count
156

**SELECT COUNT(\*) FROM etl\_project.DIM\_DATE;**

```

81 SELECT COUNT(*) FROM etl_project.DIM_LOCATION;
82 SELECT COUNT(*) FROM etl_project.DIM_ATM;
83 SELECT COUNT(*) FROM etl_project.DIM_DATE;
84 SELECT COUNT(*) FROM etl_project.DIM_CARD_TYPE;
85 SELECT COUNT(*) FROM etl_project.FACT_ATM_TRANS;

```

**Run** **Save** **Schedule** **Clear** [Send feedback](#)

**Query results** | **Table details**

Query **1427** [Execution](#) [Data](#) [Visualize](#)

Completed, started on September 08, 2021 at 22:52:23  
ELAPSED TIME: 00 m 02 s

**Rows returned (1)** [Export](#)

< 1 > [Settings](#)

count
8685

**SELECT COUNT(\*) FROM etl\_project.DIM\_CARD\_TYPE;**

```

83 SELECT COUNT(*) FROM etl_project.DIM_DATE;
84 SELECT COUNT(*) FROM etl_project.DIM_CARD_TYPE;
85 SELECT COUNT(*) FROM etl_project.FACT_ATM_TRANS;
86
87
88
89

```

**Run** **Save** **Schedule** **Clear** [Send feedback](#)

**Query results** | **Table details**

Query **1443** [Execution](#) [Data](#) [Visualize](#)

Completed, started on September 08, 2021 at 22:53:35  
ELAPSED TIME: 00 m 03 s

**Rows returned (1)** [Export](#)

< 1 > [Settings](#)

count
12

**SELECT COUNT(\*) FROM etl\_project.FACT\_ATM\_TRANS;**

83 SELECT COUNT(\*) FROM etl\_project.DIM\_DATE;  
84 SELECT COUNT(\*) FROM etl\_project.DIM\_CARD\_TYPE;  
85 SELECT COUNT(\*) FROM etl\_project.FACT\_ATM\_TRANS;  
86  
87  
88  
89

Run Save Schedule Clear

Send feedback

Query results Table details

Query 1461

Execution Data Visualize

Completed, started on September 08, 2021 at 22:54:26  
ELAPSED TIME: 00 m 02 s

Rows returned (1)

Export

Search rows

< 1 >

count

2468572